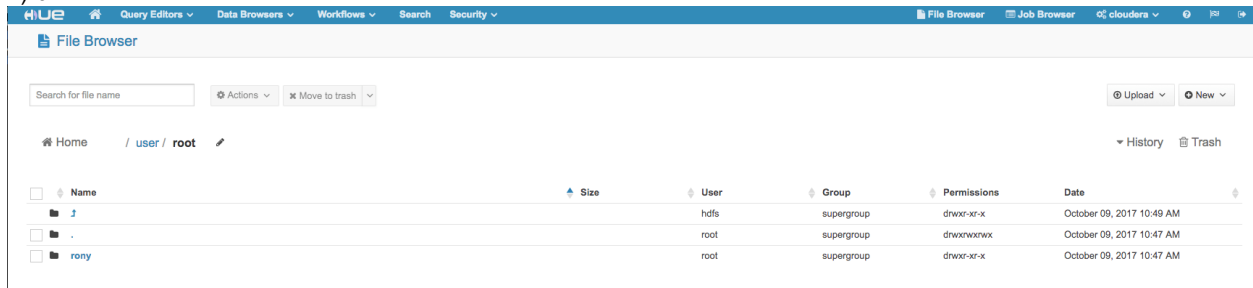


2) a.

```
[root@quickstart home]# hdfs dfs -mkdir rony
[root@quickstart home]# hdfs dfs -ls
Found 1 items
drwxr-xr-x  - root supergroup          0 2017-10-09 17:47 rony
[root@quickstart home]# _
```

2) b.



2) c.

Run instructions:

```
hdfs dfs -put bigdataclassfolder/HW1/IrishPennyJournal.txt hw1
```

```
hadoop jar /opt/cloudera/parcels/CDH-5.6.1-1.cdh5.6.1.p0.3/lib/hadoop-mapreduce/hadoop-streaming.jar -file
bigdataclassfolder/HW1/mapper.py -mapper bigdataclassfolder/HW1/mapper.py -file
bigdataclassfolder/HW1/reducer.py -reducer bigdataclassfolder/HW1/reducer.py -input hw1/ -output hw_output
```

```
hdfs dfs -get hw_output .
```

```
Launched reduce tasks=32
Data-local map tasks=2
Total time spent by all maps in occupied slots (ms)=21344
Total time spent by all reduces in occupied slots (ms)=548088
Total time spent by all map tasks (ms)=5336
Total time spent by all reduce tasks (ms)=91348
Total vcore-milliseconds taken by all map tasks=5336
Total vcore-milliseconds taken by all reduce tasks=91348
Total megabyte-milliseconds taken by all map tasks=21856256
Total megabyte-milliseconds taken by all reduce tasks=561242112
Map-Reduce Framework
  Map input records=1607
  Map output records=13394
  Map output bytes=170324
  Map output materialized bytes=128658
  Input split bytes=202
  Combine input records=0
  Combine output records=0
  Reduce input groups=9667
  Reduce shuffle bytes=128658
  Reduce input records=13394
  Reduce output records=9667
  Spilled Records=26788
  Shuffled Maps =64
  Failed Shuffles=0
  Merged Map outputs=64
  GC time elapsed (ms)=733
  CPU time spent (ms)=42310
  Physical memory (bytes) snapshot=13504847872
  Virtual memory (bytes) snapshot=85378191360
  Total committed heap usage (bytes)=54671704064
  Peak Map Physical memory (bytes)=692740096
  Peak Map Virtual memory (bytes)=2512748544
  Peak Reduce Physical memory (bytes)=392749056
  Peak Reduce Virtual memory (bytes)=2530275328
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=108605
File Output Format Counters
  Bytes Written=128415
17/10/11 15:01:38 INFO streaming.StreamJob: Output directory: hw_output
[rx294@login-1-1 ~]$ _
```

MR2

JOB ID

1504637662637_28560

TYPE

MR2

USER

rx294

STATUS

SUCCEEDED

LOGS

Logs

MAPS

100%

REDUCES

100%

DURATION

12s

streamjob5843063288978628773.jar

AttemptsTasksMetadataCounters

Text Filter

Name	Value
ID	1504637662637_28560
User	rx294
Maps	2 of 2
Reduces	32 of 32
Started	10/11/17 15:01:23
Ended	10/11/17 15:01:35
Duration	12519
Status	SUCCEEDED
dfs.balancer.address	0.0.0.0:0
dfs.balancer.block.move.timeout	0
dfs.balancer.keytab.enabled	false
dfs.balancer.max-no-move-interval	60000
dfs.block.access.key.update.interval	600
dfs.block.access.token.enable	false
dfs.block.access.token.lifetime	600
dfs.block.scanner.volume.bytes.per.second	1048576
dfs.blockreport.initialDelay	0
dfs.blockreport.intervalMsec	21600000
dfs.blockreport.split.threshold	1000000
dfs.blocksize	134217728